

Parallel Molecular Dynamics with a Time-Reversible Nose-Hoover Thermostat
on CPUs and GPUs

1.0

Generated by Doxygen 1.8.2

Mon Jan 13 2014 22:54:02

Contents

1	CBEMDGPU	1
2	Hierarchical Index	5
2.1	Class Hierarchy	5
3	Class Index	7
3.1	Class List	7
4	Class Documentation	9
4.1	atom Struct Reference	9
4.1.1	Detailed Description	9
4.2	cellList_cpu Class Reference	9
4.2.1	Detailed Description	10
4.2.2	Constructor & Destructor Documentation	10
4.2.2.1	cellList_cpu	10
4.2.3	Member Function Documentation	12
4.2.3.1	cell	12
4.2.3.2	checkUpdate	12
4.3	customException Class Reference	13
4.3.1	Detailed Description	14
4.4	float3 Struct Reference	14
4.4.1	Detailed Description	14
4.5	int3 Struct Reference	15
4.5.1	Detailed Description	15
4.6	integrator Class Reference	15
4.6.1	Detailed Description	16
4.6.2	Member Function Documentation	16
4.6.2.1	calcForce	16
4.7	nve Class Reference	17
4.7.1	Detailed Description	18
4.7.2	Member Function Documentation	18
4.7.2.1	calcForce	18

4.7.2.2	step	20
4.8	nvt_NH Class Reference	21
4.8.1	Detailed Description	22
4.8.2	Constructor & Destructor Documentation	22
4.8.2.1	nvt_NH	22
4.8.3	Member Function Documentation	22
4.8.3.1	calcForce	23
4.8.3.2	step	24
4.9	systemDefinition Class Reference	26
4.9.1	Detailed Description	27
4.9.2	Member Function Documentation	27
4.9.2.1	initRandom	27
4.9.2.2	initThermal	28
4.9.2.3	setPotential	29
4.9.2.4	writeSnapshot	29
4.10	systemProps Class Reference	30
4.10.1	Detailed Description	30
4.10.2	Member Function Documentation	30
4.10.2.1	displayCudaProps	30
4.11	SystemTest Class Reference	31
4.11.1	Detailed Description	32

Index

32

Chapter 1

CBEMDGPU

CBE MD on GPUs

by Nathan A. Mahynski, George A. Khoury, and Carmeline J. Dsilva

Our main report is located in the report/ directory (report.pdf).

Compiling

See [main.cpp](#) to set parameters which are documented by example in this file. Our code has the following dependencies:

1. Boost (<http://www.boost.org/>). Refer to Makefile for PATHTOBOOST variable which must be set. Note that compilation on TIGER is easy since you are allowed to link to other user's directories. George's installation is already provided in the Makefile.
2. Cuda toolkit 5.5 (module load module cudatoolkit/5.5.22 on TIGER)
3. Intel C++ Compilers (icpc) and OMP (module load openmpi/intel-12.1/1.4.5/64 on TIGER)
4. googletest (<https://code.google.com/p/googletest/>). Again George's GTEST_DIR directory (see Makefile) is provided to be linked against at compile time.

See FYI section below for more details. These must be taken care of before attempting to compile the code. Always use "make clean" before attempting a new compilation.

To compile the CPU version, type `$ make MD` which produces a binary called `md`

To compile the GPU version, type `$ make -f Makefile_cuda` which produces a binary called `md`

To compile tests, type `$ make TESTS` which produces a binary called `tests`

To compile the timing executable (used for CPU only scaling studies), type `$ make TIMING` which produces a binary called `timing`

To obtain the CUDA scaling results go to the `cuda_timing/` directory, type `$ make -f Makefile_cuda` which produces a binary called `cuda_timing`

To compile the program that runs a simulation to compare with LAMMPS output, type `$ make LMP_COMPARE` which produces a binary called `lmp_compare` used to generate data at specific settings we also ran on LAMMPS with.

To compile the program that tests the NVE integrator, type `$ make TEST_NVE` which produces a binary called `test_nve`

In the Makefile, the PATHTOBOOST variable should point to the directory where the C++ boost libraries are saved.

In the GTEST_DIR variable should point to the directory where the Google Tests libraries are saved.

Note that the Intel C++ compilers must be used; the GNU C++ compilers do not work with the OpenMP portion of our code (this is a known bug in the compiler, see FYI section).

Execution

Most binaries expects 4 input, the number of threads to use with OMP, the number of atoms, the skin radius for the cell/neighbor lists, and the number of steps to simulate. `$./md nthreads natoms rs nsteps > log 2> err`.

The exceptions are (1) tests which is simply executed as `./tests`, and (2) `test_nve` and (3) `Imp_compare` which are executed as `./binary_name nthreads`. However, the latter two are not of much interest; if you want to check the code is running just check to see if `./tests` works.

This also produces a `trajectory.xyz` file which can be visualized with VMD (if you have it installed) `$ vmd -xyz trajectory.xyz`

Submission scripts for TIGER are included in the `run_scaling.sh*` files. The integer suffix (i.e. `run_scaling.sh.1`) refers to the number of threads OMP will use.

Explanation of `main.cpp`

How to use, change, and make your own in 10 steps.

1. Most binaries expects 4 input, the number of threads to use with OMP, the number of atoms, the skin radius for the cell/neighbor lists, and the number of steps to simulate. `$./md nthreads natoms rs nsteps > log 2> err`

1. The random number generator seed is then set manually to ensure that results are reproducible.

1. The user must then specify basic properties about the system. A `systemDefinition` object should be instantiated and then temperature, box size, particle mass, and pair potential cutoff should be specified

```
$ systemDefinition a;
```

```
$ a.setTemp(1.0);...
```

1. The system can then be initialized with the command `initThermal` or `initRandom` (see doxygen documentation for more details). For example, in the former:

```
$ a.initThermal(nAtoms, Temp, rngSeed, separation);
```

1. Where "nAtoms" is the number of atoms, "separation" is the initial separation of the particles on the simple cubic lattice on which they are initialized, and the rest of the variables are self-explanatory. This command completely initializes the system for simulation.

1. Next, the pair potential should be specified. As an example in `main.cpp`, the preprocessor flag `NVCC` is used to select either the CPU version of the shifted lennard-jones potential (`slj`) or the GPU version (`dev_slj`). The Makefile (compare to `Makefile_cuda`) will define this if necessary and allows the code to flow naturally and work in both cases.

```
$ pointFunction_t pp = slj;
```

```
$ a.setPotential(pp);
```

1. If the GPUs are used (`NVCC` compiler flag is defined) the number of threads and blocks the GPU kernel will be invoked with must be set. This is done next.

1. Then additional arguments for the pair potential must be specified. In the case of the slj function, variables epsilon, sigma, delta, and ushift must be given. This is then set in the [systemDefinition](#).

```
$ std::vector <float> args(5);
$ args[0] = 1.0; // epsilon
$ args[1] = 1.0; // sigma
$ args[2] = 0.0; // delta
$ args[3] = 0.0; // ushift
$ a.setPotentialArgs(args);
```

1. Afterwards, the integrator (ensemble) must be specified. In the case of the NVT ensemble where we are using the Nose-Hoover thermostat, the integrator is given by the object [nvt_NH](#). This requires a damping constant, which for the slj potential should be about unity. Then the numerical timestep should be given, for the slj this is usually efficient around 0.005.

```
$ nvt_NH integrate (1.0);
$ integrate.setTimestep(timestep);
```

1. Finally the simulation is ready to iterate. A simple loop can be set to do this. An example for the case of the NVE ensemble is also provided in [test_nve.cpp](#) which can be compiled with make TEST_NVE (see [test_nve.cpp](#))

FYI

Known bugs, etc.

There are a few known instances of bugs related to compiler options, etc.

1. Using icpc instead of g++ Our code uses OMP to parallelize many calculations. Specifically, the atoms member (vector) in the [systemDefinition](#) object must be shared often. However, because it is a member of a class g++ struggles to properly share this in memory. This is a known bug in the g++ compiler which the intel (icpc) version handles rigorously. As a result, our code will produce errors if compiled with the g++ compiler when more than 1 core is used. To use this on tiger the module openmpi/intel-12.1/1.4.5/64 should be loaded to make the icpc compiler accessible. This also provides the OMP libraries necessary.
1. Cuda toolkit 5.5 CUDA is a finicky tool. Different GPUs require different toolkits and versions to work properly. In fact, compilation may succeed with a bad version but the run time behavior produces unexpected (incorrect) results. For the K20 cards on tiger, the latest toolkit (v5.5) must be loaded. To do so, load the module cudatoolkit/5.5.22 before attempting to compile the program. The Makefile must include flags consistent with the GPUs version of CUDA (which on TIGER is 3.5) so the Makefile_cuda contains a flag "NVFLAGS = -generate arch=compute_35,code=sm_35" for the .cu files. Furthermore, you will find that preprocessor flags NVCC and NOGPU are found throughout the code which act as switches to activate/deactivate GPU functionality throughout the compilation process.
1. Tiger GPUs Unfortunately it appears that the queueing system on tiger is having problems reserving all or some of the GPUs exclusively for single jobs by users. As a result, thrust (the GPU equivalent of the STL for C++) will have memory issues if one or more of the GPUs on a node are already in use. Because of the unusually high load on tiger over the past month, we have been unable to obtain good results on this cluster since usually this situation is encountered. Our private GPU cluster was used to obtain our results instead, though our Makefile_cuda is set so this should compile properly on tiger if you want to check.
1. Use of preprocessor flags The use of NVCC and NOGPU flags is crucial during compilation to ensure the correct header files are linked and compiled, and certain sections of code are ignored depending on the compiler. The Makefiles provided already handle this appropriately, but one should be aware of these details while reading the code.

Chapter 2

Hierarchical Index

2.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

atom	9
cellList_cpu	9
exception	
customException	13
float3	14
int3	15
integrator	15
nve	17
nvt_NH	21
systemDefinition	26
systemProps	30
Test	
SystemTest	31

Chapter 3

Class Index

3.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

atom	Structure of an atom	9
cellList_cpu	9
customException	Error reporting	13
float3	3 floating point numbers, same as defined for GPUs	14
int3	3 integers, same as defined for GPUs	15
integrator	Base class for integrators such as NVT (Nose-Hoover) or NVE ensembles	15
nve	Integration scheme that preserves total energy of the system	17
nvt_NH	Uses Nose-Hoover integration method to thermostat a system (constant T rather than E)	21
systemDefinition	Contains all information pertaining to a system being simulated	26
systemProps	Holds all information about the CPU and GPU the simulation is being performed on	30
SystemTest	31

Chapter 4

Class Documentation

4.1 atom Struct Reference

Structure of an atom.

```
#include <dataTypes.h>
```

Public Attributes

- [float3 pos](#)
Position.
- [float3 vel](#)
Velocity.
- [float3 acc](#)
Acceleration.

4.1.1 Detailed Description

Structure of an atom.

Definition at line 29 of file `dataTypes.h`.

The documentation for this struct was generated from the following file:

- `/Users/nathanmahynski/CBEMDGPU/dataTypes.h`

4.2 cellList_cpu Class Reference

```
#include <cellList.h>
```

Public Member Functions

- [cellList_cpu](#) (const [float3](#) &box, const float rc, const float rs)
- void [checkUpdate](#) (const [systemDefinition](#) &sys)
Check if the neighbor list requires updating.
- int [cell](#) (const [float3](#) &pos)
Calculate the cell in which a given coordinate is located.
- int [head](#) (const int [cell](#)) const

Return the first atom (aka 'head') of each cell.

- `int list` (const int index) const

Iterates through a linked list of particles, returns the index of the next atom in line.

- `std::vector< int > neighbors` (const int cellID) const

Returns the indices of a cell's neighboring cells.

Public Attributes

- `int3 nCells`

Number of cells in each direction.

4.2.1 Detailed Description

Cell Lists

Author

Nathan A. Mahynski

Date

11/19/13

Maintains a linked list to track cells on the CPU

Definition at line 39 of file cellList.h.

4.2.2 Constructor & Destructor Documentation

4.2.2.1 `cellList_cpu::cellList_cpu (const float3 & box, const float rc, const float rs)`

Cell Lists

Author

Nathan A. Mahynski

Date

11/19/13

Initialize a cell list

Parameters

<code>in</code>	<code>box</code>	Box size
<code>in</code>	<code>rc</code>	Cutoff radius
<code>in</code>	<code>rs</code>	Skin Radius

Definition at line 156 of file cellList.cpp.

References `nCells`.

```

if (rc < 0.0) {
    throw customException("Cutoff radius must be > 0");
    return;
}

```

```

    }
    rc_ = rc;
    if (rs < 0.0) {
        throw customException("Skin radius must be > 0");
        return;
    }
    rs_ = rs;

    box_ = box;

    start_ = 1;
    lcell_.x = 1.01*(rc+rs);
    nCells.x = (int) floor (box.x/lcell_.x);
    lcell_.x = (box.x/nCells.x);

    lcell_.y = 1.01*(rc+rs);
    nCells.y = (int) floor (box.y/lcell_.y);
    lcell_.y = (box.y/nCells.y);

    lcell_.z = 1.01*(rc+rs);
    nCells.z = (int) floor (box.z/lcell_.z);
    lcell_.z = (box.z/nCells.z);

    if (lcell_.x <= (rc+rs) || lcell_.y <= (rc+rs) || lcell_.z < (rc+rs)) {
        throw customException("Cell width must exceed sum of cutoff
            and skin radius");
        return;
    }

    if (lcell_.x < 1) {
        throw customException ("Box dimension x too small relative
            to skin and cutoff radius to use cell lists");
        return;
    }
    if (lcell_.y < 1) {
        throw customException ("Box dimension y too small relative
            to skin and cutoff radius to use cell lists");
        return;
    }
    if (lcell_.z < 1) {
        throw customException ("Box dimension z too small relative
            to skin and cutoff radius to use cell lists");
        return;
    }

    if (nCells.x < 3 || nCells.y < 3 || nCells.z < 3) {
        throw customException("Must be able to have at least 3 cells
            in each direction, change box size, skin, or cutoff radius");
        return;
    }

    try {
        head_.resize(nCells.x*nCells.y*nCells.z, -1);
    } catch (std::exception &e) {
        std::cerr << e.what() << std::endl;
        throw customException ("Unable to initialize head for cell
            list");
        return;
    }
    // build neighbors for each cell
    neighbor_.resize(nCells.x*nCells.y*nCells.z);
    for (unsigned int cellID = 0; cellID < nCells.x*nCells.y*nCells
        .z; ++cellID) {
        const int zref = floor(cellID/(nCells.x*nCells.y));
        const int yref = floor((cellID - zref*(nCells.x*nCells.y))/
            nCells.y);
        const int xref = floor(cellID - zref*(nCells.x*nCells.y) -
            yref*nCells.x);
        for (int dx = -1; dx <= 1; ++dx) {
            int xcell = xref + dx;
            if (xcell >= nCells.x) xcell = 0;
            if (xcell < 0) xcell = nCells.x-1;
            for (int dy = -1; dy <= 1; ++dy) {
                int ycell = yref + dy;
                if (ycell >= nCells.y) ycell = 0;
                if (ycell < 0) ycell = nCells.y-1;
                for (int dz = -1; dz <= 1; ++dz) {
                    int zcell = zref + dz;
                    if (zcell >= nCells.z) zcell = 0;
                    if (zcell < 0) zcell = nCells.z-1;
                    const int cellID2 = xcell + ycell*nCells.x + zcell*(nCells
                        .x*nCells.y);
                    neighbor_[cellID].push_back(cellID2);
                }
            }
        }
    }
}

```

```

    for (unsigned int cellID = 0; cellID < nCells.x*nCells.y*nCells
        .z; ++cellID) {
    if (neighbor_[cellID].size() != 27) {
        throw customException ("Cell list initial build failed
            to find 27 neighbors (including self)");
        return;
    }
    }
}

```

4.2.3 Member Function Documentation

4.2.3.1 int cellList_cpu::cell (const float3 & pos)

Calculate the cell in which a given coordinate is located.

Calculates the cell (linear index of it) a position belongs to in a periodic box. The position does not need to be "inside" the box, boundary conditions are applied.

Parameters

in	pos	Position of atom
----	-----	------------------

Returns

cell

Definition at line 252 of file cellList.cpp.

References nCells.

Referenced by checkUpdate(), and head().

```

{
    float3 inBox = pbc(pos, box_);
    const int x = (int) floor(inBox.x/lcell_.x);
    const int y = (int) floor(inBox.y/lcell_.y);
    const int z = (int) floor(inBox.z/lcell_.z);
    return x + y*nCells.x + z*nCells.x*nCells.y;
}

```

4.2.3.2 void cellList_cpu::checkUpdate (const systemDefinition & sys)

Check if the neighbor list requires updating.

Checks to see if the cell list needs to be updated and if so, rebuilds Uses linked list, so algorithm in O(N)

Parameters

in	sys	System definition
----	-----	-------------------

Definition at line 266 of file cellList.cpp.

References systemDefinition::atoms, systemDefinition::box(), cell(), and systemDefinition::numAtoms().

Referenced by integrator::calcForce().

```

{
    int build = 0;
    if (start_) {
        const int natoms = sys.numAtoms();
        try {
            list_.resize(natoms, -1);
        } catch (std::exception& e) {
            std::cerr << e.what() << std::endl;
            throw customException ("Unable to initialize list
                for cell list");
        }
    }
}

```



```

        return;
    }
    try {
        posAtLastBuild_.resize(natoms);
    } catch (std::exception& e) {
        std::cerr << e.what() << std::endl;
        throw customException ("Unable to initialize
position list for cell list");
        return;
    }
    start_ = 0;
    build = 1;
} else {
    // check max displacements
    float drMax1_ = 0.0;
    float drMax2_ = 0.0;
    float3 dummy;
    for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
        const float dr2 = pbcDist2 (sys.atoms[i].pos,
posAtLastBuild_[i], dummy, sys.box());
        if (dr2 > drMax1_*drMax1_) {
            drMax2_ = drMax1_;
            drMax1_ = sqrt(dr2);
        } else if (dr2 > drMax2_*drMax2_) {
            drMax2_ = sqrt(dr2);
        }
    }
    if (drMax1_+drMax2_ > rs_) {
        build = 1;
    } else {
        build = 0;
    }
}

if (build) {
    if (sys.numAtoms() != list_.size()) {
        throw customException ("Number of atoms in
simulation has changed");
        return;
    }
    for (unsigned int i = 0; i < head_.size(); ++i) {
        head_[i] = -1;
    }
    for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
        const int icell = cell(sys.atoms[i].pos);
        list_[i] = head_[icell];
        head_[icell] = i;
        posAtLastBuild_[i] = sys.atoms[i].pos;
    }
}

return;
}

```

The documentation for this class was generated from the following files:


- /Users/nathanmahynski/CBEMDGPU/cellList.h
- /Users/nathanmahynski/CBEMDGPU/cellList.cpp

4.3 customException Class Reference

Error reporting.

```
#include <common.h>
```

Inheritance diagram for customException:



Public Member Functions

- `const char * what () const throw ()`
Return a message pertaining to what went wrong.
- `customException (std::string m="custom exception occurred")`
Instantiate the object with a user specified message.

4.3.1 Detailed Description

Error reporting.

Common values, structures, etc.

Author

Nathan A. Mahynski

Date

11/17/13

Definition at line 14 of file `common.h`.

The documentation for this class was generated from the following file:

- `/Users/nathanmahynski/CBEMDGPU/common.h`

4.4 float3 Struct Reference

3 floating point numbers, same as defined for GPUs

```
#include <dataTypes.h>
```

Public Attributes

- `float x`
- `float y`
- `float z`

4.4.1 Detailed Description

3 floating point numbers, same as defined for GPUs

Requisite 'optimal' data types.

Author

Nathan A. Mahynski

Date

11/21/13

Definition at line 16 of file `dataTypes.h`.

The documentation for this struct was generated from the following file:

- `/Users/nathanmahynski/CBEMDGPU/dataTypes.h`

4.5 int3 Struct Reference

3 integers, same as defined for GPUs

```
#include <dataTypes.h>
```

Public Attributes

- int **x**
- int **y**
- int **z**

4.5.1 Detailed Description

3 integers, same as defined for GPUs

Definition at line 22 of file `dataTypes.h`.

The documentation for this struct was generated from the following file:


- `/Users/nathanmahynski/CBEMDGPU/dataTypes.h`

4.6 integrator Class Reference

Base class for integrators such as NVT (Nose-Hoover) or NVE ensembles.

```
#include <integrator.h>
```

Inheritance diagram for integrator:



Public Member Functions

- void [setTimestep](#) (const float dt)
Set the integrator timestep.
- void [calcForce](#) ([systemDefinition](#) &sys)
Calculate the forces on each atom.
- virtual void [step](#) ([systemDefinition](#) &sys)=0
Move the system forward a step in time.

Protected Attributes

- [cellList_cpu](#) [cl_](#)
Cell or neighbor list.
- `std::vector< float3 >` [lastAccelerations_](#)
Acceleration of particles on previous timestep (useful for NVE integrator)
- float [dt_](#)
Timestep size.

- int [start_](#)

Flag for whether this object has been initialized or not.

4.6.1 Detailed Description

Base class for integrators such as NVT (Nose-Hoover) or NVE ensembles.

Integrator

Author

Nathan A. Mahynski

Date

11/17/13

Definition at line 15 of file integrator.h.

4.6.2 Member Function Documentation

4.6.2.1 void integrator::calcForce (systemDefinition & sys)

Calculate the forces on each atom.

Integration

Author

Nathan A. Mahynski

Date

11/19/13

Calculate the pairwise forces in a system. This also calculates the potential energy of a system. The kinetic energy is calculated during the verlet integration.

Parameters

<code>in, out</code>	<code>sys</code>	System definition
----------------------	------------------	-------------------

Calculate the pairwise forces in a system. This also calculates the potential energy of a system. The kinetic energy is calculated during the verlet integration.

Parameters

<code>in, out</code>	<code>sys</code>	System definition
----------------------	------------------	-------------------

Definition at line 23 of file integrator.cpp.

References `systemDefinition::atoms`, `systemDefinition::box()`, `cellList_cpu::checkUpdate()`, `cl_`, `cellList_cpu::head()`, `cellList_cpu::list()`, `systemDefinition::mass()`, `cellList_cpu::nCells`, `cellList_cpu::neighbors()`, `systemDefinition::numAtoms()`, `systemDefinition::potential`, `systemDefinition::potentialArgs()`, `systemDefinition::rcut()`, and `systemDefinition::setPotE()`.

Referenced by `nve::step()`, and `nvt_NH::step()`.

```

    {
// For cache coherency allocate new space for calculations
float3 empty;
empty.x = 0; empty.y = 0; empty.z = 0;
std::vector<float3> acc (sys.numAtoms(), empty);
const float rc = sys.rcut();

// every time, check if the cell list needs to be updated first
cl_.checkUpdate(sys);

float Up = 0.0;
const float3 box = sys.box();
const float invMass = 1.0/sys.mass();

// traverse cell list and calculate total system potential energy
std::vector<float> args = sys.potentialArgs();
#pragma omp parallel for reduction(+:Up) schedule(dynamic, 1)
for (unsigned int cellID = 0; cellID < cl_.nCells.x*cl_.nCells
.y*cl_.nCells.z; ++cellID) {
    int atom1 = cl_.head(cellID);
    while (atom1 >= 0) {
        std::vector<int> neighbors = cl_.neighbors(cellID);
        for (int index = 0; index < neighbors.size(); ++index) {
            const int cellID2 = neighbors[index];
            int atom2 = cl_.head(cellID2);
            while (atom2 >= 0) {
                if (atom1 > atom2) {
                    float3 pf;
                    Up += sys.potential (&sys.atoms[atom1].
pos, &sys.atoms[atom2].pos, &pf, &box, &args[0], &rc);
                    acc[atom1].x -= pf.x*invMass;
                    acc[atom1].y -= pf.y*invMass;
                    acc[atom1].z -= pf.z*invMass;
                    acc[atom2].x += pf.x*invMass;
                    acc[atom2].y += pf.y*invMass;
                    acc[atom2].z += pf.z*invMass;
                }
                atom2 = cl_.list(atom2);
            }
        }
        atom1 = cl_.list(atom1);
    }
}

// save acceleration in array of atoms in system
#pragma omp parallel for
for (int i = 0; i < sys.numAtoms(); ++i) {
    sys.atoms[i].acc.x = -acc[i].x;
    sys.atoms[i].acc.y = -acc[i].y;
    sys.atoms[i].acc.z = -acc[i].z;
}

// set Up
sys.setPotE(Up);
}

```

The documentation for this class was generated from the following files:


- /Users/nathanmahynski/CBEMDGPU/integrator.h
- /Users/nathanmahynski/CBEMDGPU/integrator.cpp
- /Users/nathanmahynski/CBEMDGPU/integrator.cu

4.7 nve Class Reference

Integration scheme that preserves total energy of the system.

```
#include <nve.h>
```

Inheritance diagram for nve:



Public Member Functions

- void [step](#) ([systemDefinition](#) &sys)
- void [setTimestep](#) (const float dt)
Set the integrator timestep.
- void [calcForce](#) ([systemDefinition](#) &sys)
Calculate the forces on each atom.

Protected Attributes

- [cellList_cpu](#) [cl_](#)
Cell or neighbor list.
- `std::vector< float3 > lastAccelerations_`
Acceleration of particles on previous timestep (useful for NVE integrator)
- float [dt_](#)
Timestep size.
- int [start_](#)
Flag for whether this object has been initialized or not.

4.7.1 Detailed Description

Integration scheme that preserves total energy of the system.

NVE integration.

Author

Nathan A. Mahynski

Date

11/18/13

Definition at line 14 of file nve.h.

4.7.2 Member Function Documentation

4.7.2.1 void [integrator::calcForce](#) ([systemDefinition](#) & sys) `[inherited]`

Calculate the forces on each atom.

Integration

Author

Nathan A. Mahynski

Date

11/19/13

Calculate the pairwise forces in a system. This also calculates the potential energy of a system. The kinetic energy is calculated during the verlet integration.

Parameters

in, out	sys	System definition
---------	-----	-------------------

Calculate the pairwise forces in a system. This also calculates the potential energy of a system. The kinetic energy is calculated during the verlet integration.

Parameters

in, out	sys	System definition
---------	-----	-------------------

Definition at line 23 of file integrator.cpp.

References `systemDefinition::atoms`, `systemDefinition::box()`, `cellList_cpu::checkUpdate()`, `integrator::cl_`, `cellList_cpu::head()`, `cellList_cpu::list()`, `systemDefinition::mass()`, `cellList_cpu::nCells`, `cellList_cpu::neighbors()`, `systemDefinition::numAtoms()`, `systemDefinition::potential`, `systemDefinition::potentialArgs()`, `systemDefinition::rcut()`, and `systemDefinition::setPotE()`.

Referenced by `step()`, and `nvt_NH::step()`.

```

{
    // For cache coherency allocate new space for calculations
    float3 empty;
    empty.x = 0; empty.y = 0; empty.z = 0;
    std::vector<float3> acc (sys.numAtoms(), empty);
    const float rc = sys.rcut();

    // every time, check if the cell list needs to be updated first
    cl_.checkUpdate(sys);

    float Up = 0.0;
    const float3 box = sys.box();
    const float invMass = 1.0/sys.mass();

    // traverse cell list and calculate total system potential energy
    std::vector<float> args = sys.potentialArgs();
    #pragma omp parallel for reduction(+:Up) schedule(dynamic, 1)
    for (unsigned int cellID = 0; cellID < cl_.nCells.x*cl_.nCells
        .y*cl_.nCells.z; ++cellID) {
        int atom1 = cl_.head(cellID);
        while (atom1 >= 0) {
            std::vector<int> neighbors = cl_.neighbors(cellID);
            for (int index = 0; index < neighbors.size(); ++index) {
                const int cellID2 = neighbors[index];
                int atom2 = cl_.head(cellID2);
                while (atom2 >= 0) {
                    if (atom1 > atom2) {
                        float3 pf;
                        Up += sys.potential (&sys.atoms[atom1].
pos, &sys.atoms[atom2].pos, &pf, &box, &args[0], &rc);
                        acc[atom1].x -= pf.x*invMass;
                        acc[atom1].y -= pf.y*invMass;
                        acc[atom1].z -= pf.z*invMass;
                        acc[atom2].x += pf.x*invMass;
                        acc[atom2].y += pf.y*invMass;
                        acc[atom2].z += pf.z*invMass;
                    }
                    atom2 = cl_.list(atom2);
                }
            }
            atom1 = cl_.list(atom1);
        }
    }

    // save acceleration in array of atoms in system
    #pragma omp parallel for
    for (int i = 0; i < sys.numAtoms(); ++i) {
        sys.atoms[i].acc.x = -acc[i].x;
        sys.atoms[i].acc.y = -acc[i].y;
        sys.atoms[i].acc.z = -acc[i].z;
    }

    // set Up
    sys.setPotE(Up);
}

```

4.7.2.2 void nve::step (systemDefinition & sys) [virtual]

Do NVE integration

Author

Nathan A. Mahynski

Date

11/18/13

Integrate a single timestep forward using Velocity-Verlet integration scheme. Update the positions and velocities using velocity verlet integration. This uses the accelerations stored on each atom. Creates a cell list the first time it is called.

Parameters

in, out	sys	System definition
---------	-----	-------------------

Implements [integrator](#).

Definition at line 23 of file nve.cpp.

References [systemDefinition::atoms](#), [systemDefinition::box\(\)](#), [integrator::calcForce\(\)](#), [integrator::cl_](#), [integrator::dt_](#), [systemDefinition::mass\(\)](#), [systemDefinition::numAtoms\(\)](#), [systemDefinition::rcut\(\)](#), [systemDefinition::rskin\(\)](#), [systemDefinition::setKinE\(\)](#), [integrator::start_](#), and [systemDefinition::updateInstantTemp\(\)](#).

```

{
    int chunk = OMP_CHUNK;
    if (start_) {
        try {
            cellList_cpu tmpCL (sys.box(), sys.rcut(), sys.
rskin());
            cl_ = tmpCL;
        } catch (std::exception &e) {
            std::cerr << e.what() << std::endl;
            throw customException("Failed to integrate on first
step");
        }

        // get initial temperature
        calcForce(sys);
        float tmp=0.0, Uk = 0.0;
        #pragma omp parallel for reduction(+:Uk)
        for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
            Uk += (sys.atoms[i].vel.x*sys.atoms[i].vel.x)+(sys.atoms
[i].vel.y*sys.atoms[i].vel.y)+(sys.atoms[i].vel.z*sys.atoms[i].
vel.z);
        }
        Uk *= sys.mass();
        tmp = Uk;
        Uk *= 0.5;
        tmp /= (3.0*(sys.numAtoms()-1.0));
        sys.updateInstantTemp(tmp);
        sys.setKinE(Uk);
        start_ = 0;
    }

    // (1) evolve particle velocities
    #pragma omp parallel shared(sys)
    {
        #pragma omp for schedule(dynamic,OMP_CHUNK)
        for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
            sys.atoms[i].vel.x += 0.5*dt_*(sys.atoms[i].acc.x);
            sys.atoms[i].vel.y += 0.5*dt_*(sys.atoms[i].acc.y);
            sys.atoms[i].vel.z += 0.5*dt_*(sys.atoms[i].acc.z);
        }

        // (2) evolve particle positions
        #pragma omp for schedule(dynamic,OMP_CHUNK)
        for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
            sys.atoms[i].pos.x += sys.atoms[i].vel.x*dt_;
            sys.atoms[i].pos.y += sys.atoms[i].vel.y*dt_;
            sys.atoms[i].pos.z += sys.atoms[i].vel.z*dt_;
        }
    }
}

```



```

    }
}

// (3) calc force
calcForce(sys);

// (4) evolve particle velocities
#pragma omp parallel shared(sys)
{
    #pragma omp for schedule(dynamic,OMP_CHUNK)
    for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
        sys.atoms[i].vel.x += 0.5*dt_*(sys.atoms[i].acc.x);
        sys.atoms[i].vel.y += 0.5*dt_*(sys.atoms[i].acc.y);
        sys.atoms[i].vel.z += 0.5*dt_*(sys.atoms[i].acc.z);
    }
}
float Uk = 0.0;
float tmp = 0.0;

#pragma omp parallel for reduction(+:Uk)
// get temperature and kinetic energy
for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
    Uk += (sys.atoms[i].vel.x*sys.atoms[i].vel.x)+(sys.atoms
[i].vel.y*sys.atoms[i].vel.y)+(sys.atoms[i].vel.z*sys.atoms[i].
vel.z);
}

Uk *= sys.mass();
tmp = Uk;
Uk *= 0.5;
tmp /= (3.0*(sys.numAtoms()-1.0));
sys.updateInstantTemp(tmp);
sys.setKinE(Uk);
}

```

The documentation for this class was generated from the following files:


- /Users/nathanmahynski/CBEMDGPU/nve.h
- /Users/nathanmahynski/CBEMDGPU/nve.cpp

4.8 nvt_NH Class Reference

Uses Nose-Hoover integration method to thermostat a system (constant T rather than E)

```
#include <nvt.h>
```

Inheritance diagram for nvt_NH:



Public Member Functions

- [nvt_NH](#) (const float Q)
- void [step](#) (systemDefinition &sys)
- void [setTimestep](#) (const float dt)
 - Set the integrator timestep.*
- void [calcForce](#) (systemDefinition &sys)
 - Calculate the forces on each atom.*

Protected Attributes

- [cellList_cpu cl_](#)
Cell or neighbor list.
- `std::vector< float3 > lastAccelerations_`
Acceleration of particles on previous timestep (useful for NVE integrator)
- `float dt_`
Timestep size.
- `int start_`
Flag for whether this object has been initialized or not.

4.8.1 Detailed Description

Uses Nose-Hover integration method to thermostat a system (constant T rather than E)
NVT integration with Nose-Hoover thermostat.

Author

Nathan A. Mahynski

Date

11/18/13

Definition at line 14 of file nvt.h.

4.8.2 Constructor & Destructor Documentation

4.8.2.1 `nvt_NH::nvt_NH (const float Q)`

Do NVT integration with Nose-Hoover thermostat.

Date

11/18/13

Initialize integrator

Parameters

<code>in</code>	<code>Q</code>	Thermal mass for thermostat
-----------------	----------------	-----------------------------

Definition at line 20 of file nvt.cpp.

References `integrator::start_`.

```

{
    Q_ = Q;
    gamma_ = 0.0;
    start_ = 1;
}
```

4.8.3 Member Function Documentation

4.8.3.1 void integrator::calcForce (systemDefinition & sys) [inherited]

Calculate the forces on each atom.

Integration

Author

Nathan A. Mahynski

Date

11/19/13

Calculate the pairwise forces in a system. This also calculates the potential energy of a system. The kinetic energy is calculated during the verlet integration.

Parameters

in, out	sys	System definition
---------	-----	-------------------

Calculate the pairwise forces in a system. This also calculates the potential energy of a system. The kinetic energy is calculated during the verlet integration.

Parameters

in, out	sys	System definition
---------	-----	-------------------

Definition at line 23 of file integrator.cpp.

References systemDefinition::atoms, systemDefinition::box(), cellList_cpu::checkUpdate(), integrator::cl_, cellList_cpu::head(), cellList_cpu::list(), systemDefinition::mass(), cellList_cpu::nCells, cellList_cpu::neighbors(), systemDefinition::numAtoms(), systemDefinition::potential, systemDefinition::potentialArgs(), systemDefinition::rcut(), and systemDefinition::setPotE().

Referenced by nve::step(), and step().

```

{
    // For cache coeherency allocate new space for calculations
    float3 empty;
    empty.x = 0; empty.y = 0; empty.z = 0;
    std::vector<float3> acc (sys.numAtoms(), empty);
    const float rc = sys.rcut();

    // every time, check if the cell list needs to be updated first
    cl_.checkUpdate(sys);

    float Up = 0.0;
    const float3 box = sys.box();
    const float invMass = 1.0/sys.mass();

    // traverse cell list and calculate total system potential energy
    std::vector<float> args = sys.potentialArgs();
#pragma omp parallel for reduction(+:Up) schedule(dynamic, 1)
    for (unsigned int cellID = 0; cellID < cl_.nCells.x*cl_.nCells.y*cl_.nCells.z; ++cellID) {
        int atom1 = cl_.head(cellID);
        while (atom1 >= 0) {
            std::vector<int> neighbors = cl_.neighbors(cellID);
            for (int index = 0; index < neighbors.size(); ++index) {
                const int cellID2 = neighbors[index];
                int atom2 = cl_.head(cellID2);
                while (atom2 >= 0) {
                    if (atom1 > atom2) {
                        float3 pf;
                        Up += sys.potential (&sys.atoms[atom1].pos, &sys.atoms[atom2].pos, &pf, &box, &args[0], &rc);
                        acc[atom1].x -= pf.x*invMass;
                        acc[atom1].y -= pf.y*invMass;
                        acc[atom1].z -= pf.z*invMass;
                        acc[atom2].x += pf.x*invMass;
                    }
                }
            }
        }
    }
}

```

```

        acc[atom2].y += pf.y*invMass;
        acc[atom2].z += pf.z*invMass;
    }
    atom2 = cl_.list(atom2);
}
}
atom1 = cl_.list(atom1);
}
}

// save acceleration in array of atoms in system
#pragma omp parallel for
for (int i = 0; i < sys.numAtoms(); ++i) {
    sys.atoms[i].acc.x = -acc[i].x;
    sys.atoms[i].acc.y = -acc[i].y;
    sys.atoms[i].acc.z = -acc[i].z;
}

// set Up
sys.setPotE(Up);
}

```

4.8.3.2 void nvt.NH::step (systemDefinition & sys) [virtual]

Integrate a single timestep forward using Velocity-Verlet integration scheme. This is NOT identical since some intermediate bookkeeping needs to be handled for thermostat. Creates a cell list the first time it is called.

Parameters

in, out	sys	System definition
---------	-----	-------------------

Implements [integrator](#).

Definition at line 32 of file nvt.cpp.

References [systemDefinition::atoms](#), [systemDefinition::box\(\)](#), [integrator::calcForce\(\)](#), [integrator::cl_](#), [integrator::dt_](#), [systemDefinition::instantT\(\)](#), [systemDefinition::mass\(\)](#), [systemDefinition::numAtoms\(\)](#), [systemDefinition::rcut\(\)](#), [systemDefinition::rskin\(\)](#), [systemDefinition::setKinE\(\)](#), [integrator::start_](#), [systemDefinition::targetT\(\)](#), and [systemDefinition::updateInstantTemp\(\)](#).

```

{
    int chunk = OMP_CHUNK;
    if (start_) {
        try {
            cellList_cpu tmpCL (sys.box(), sys.rcut(), sys.
rskin());
            cl_ = tmpCL;
        } catch (std::exception &e) {
            std::cerr << e.what() << std::endl;
            throw customException("Failed to integrate on first
step");
        }

        gamma_ = 0.0;
        for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
            gamma_ += (sys.atoms[i].vel.x*sys.atoms[i].vel.x)+(sys.
atoms[i].vel.y*sys.atoms[i].vel.y)+(sys.atoms[i].vel.z*sys.atoms
[i].vel.z);
        }
        gamma_ -= (3.0*(sys.numAtoms()-1.0))*sys.instantT();
        gamma_ /= Q_;

        // get initial temperature
        calcForce(sys);

        float tmp=0.0, Uk = 0.0;
        #pragma omp parallel for reduction(+:Uk)
        for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
            Uk += (sys.atoms[i].vel.x*sys.atoms[i].vel.x)+(sys.atoms
[i].vel.y*sys.atoms[i].vel.y)+(sys.atoms[i].vel.z*sys.atoms[i].
vel.z);
        }
        Uk *= sys.mass();
        tmp = Uk;
        Uk *= 0.5;
        tmp /= (3.0*(sys.numAtoms()-1.0));
        sys.updateInstantTemp(tmp);
    }
}

```

```

        sys.setKinE(Uk);
        start_ = 0;
        gammadot_ = 0.0;
        gammadd_ = 0.0;
    }

    tau2_ = Q_ / ((3.0*(sys.numAtoms()-1.0))*sys.targetT());
    gammadd_ = 1/tau2_*(sys.instantT()/sys.targetT()-1);

    // (1) update thermostat velocity and thermostat position
    // velocity half-step
    gammadot_ += dt_*0.5*gammadd_;

    // position step
    gamma_ += gammadot_*dt_;

    // (2) evolve particle velocities
    #pragma omp parallel shared(sys)
    {
        #pragma omp for schedule(dynamic,OMP_CHUNK)
        for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
            sys.atoms[i].vel.x = sys.atoms[i].vel.x*exp(-gammadot_*
            dt_*0.5) + 0.5*dt_*(sys.atoms[i].acc.x);
            sys.atoms[i].vel.y = sys.atoms[i].vel.y*exp(-gammadot_*
            dt_*0.5) + 0.5*dt_*(sys.atoms[i].acc.y);
            sys.atoms[i].vel.z = sys.atoms[i].vel.z*exp(-gammadot_*
            dt_*0.5) + 0.5*dt_*(sys.atoms[i].acc.z);
        }

        // (3) evolve particle positions
        #pragma omp for schedule(dynamic,OMP_CHUNK)
        for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
            sys.atoms[i].pos.x += sys.atoms[i].vel.x*dt_;
            sys.atoms[i].pos.y += sys.atoms[i].vel.y*dt_;
            sys.atoms[i].pos.z += sys.atoms[i].vel.z*dt_;
        }
    }

    // (4) calc force
    calcForce(sys);

    // (5) evolve particle velocities
    #pragma omp parallel shared(sys)
    {
        #pragma omp for schedule(dynamic,OMP_CHUNK)
        for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
            sys.atoms[i].vel.x = (sys.atoms[i].vel.x+sys.atoms[i].
            acc.x*dt_*0.5)*exp(-gammadot_*dt_*0.5);
            sys.atoms[i].vel.y = (sys.atoms[i].vel.y+sys.atoms[i].
            acc.y*dt_*0.5)*exp(-gammadot_*dt_*0.5);
            sys.atoms[i].vel.z = (sys.atoms[i].vel.z+sys.atoms[i].
            acc.z*dt_*0.5)*exp(-gammadot_*dt_*0.5);
        }
    }

    float Uk = 0.0;
    float tmp = 0.0;

    #pragma omp parallel for reduction(+:Uk)
    // get temperature and kinetic energy
    for (unsigned int i = 0; i < sys.numAtoms(); ++i) {
        Uk += (sys.atoms[i].vel.x*sys.atoms[i].vel.x)+(sys.atoms
        [i].vel.y*sys.atoms[i].vel.y)+(sys.atoms[i].vel.z*sys.atoms[i].
        vel.z);
    }

    Uk *= sys.mass();
    tmp = Uk;
    Uk *= 0.5;
    tmp /= (3.0*(sys.numAtoms()-1.0));
    sys.updateInstantTemp(tmp);
    sys.setKinE(Uk);

    // (6) update thermostat velocity
    gammadd_ = 1/tau2_*(sys.instantT()/sys.targetT()-1);
    gammadot_ += dt_*0.5*gammadd_;
}

```

The documentation for this class was generated from the following files:

- /Users/nathanmahynski/CBEMDGPU/nvt.h
- /Users/nathanmahynski/CBEMDGPU/nvt.cpp

4.9 systemDefinition Class Reference

Contains all information pertaining to a system being simulated.

```
#include <system.h>
```

Public Member Functions

- void [initRandom](#) (const int N, const int rngSeed)
- void [initThermal](#) (const int N, const float Tset, const int rngSeed, const float dx)
- void [updateInstantTemp](#) (const float T)
Manually set the instantaneous temperature.
- void [setTemp](#) (const float T)
Assign the target temperature for NVT simulations.
- void [setMass](#) (const float m)
Assign the mass of each particle.
- void [setRcut](#) (const float rc)
Assign the cutoff radius of the pair potential.
- void [setRskin](#) (const float rs)
Assign the skin radius as a buffer for the neighbor/cell lists.
- void [setBox](#) (const float x, const float y, const float z)
Assign the simulation box size.
- void [printBox](#) ()
Print the box dimesions to stdout.
- [float3 box](#) () const
Report the box dimensions.
- [float instantT](#) () const
Report the instantaneous temperature.
- [float targetT](#) () const
Report the target temperature for NVT simulations.
- [float mass](#) () const
Report the particle's mass.
- [float PotE](#) () const
Report the instantaneous potential energy of the system.
- void [setPotE](#) (const float Up)
Assign the potential energy.
- void [setKinE](#) (const float Uk)
Assign the kinetic energy.
- [float KinE](#) () const
Report the instantaneous kinetic energy of the system.
- [float rskin](#) () const
Report the skin radius for the neighbor/cell lists.
- [float rcut](#) () const
Report the pair potential function cutoff radius.
- void [writeSnapshot](#) ()
Print a snapshot of the system.
- int [numAtoms](#) () const
Report the pair potential function cutoff radius.
- void [setPotentialArgs](#) (const std::vector< float > args)
Assign additional arguments to the pair potential function.
- std::vector< float > [potentialArgs](#) ()
Report additional arguments to the pair potential function.
- void [setPotential](#) (pointFunction_t pp)
Assign pair potential function.

Public Attributes

- pointFunction_t [potential](#)
Pointer to pair potential function.
- std::vector< [atom](#) > [atoms](#)
Vector of atoms in the system.

4.9.1 Detailed Description

Contains all information pertaining to a system being simulated.

System definition

Author

Nathan A. Mahynski

Date

11/17/13

Definition at line 17 of file system.h.

4.9.2 Member Function Documentation

4.9.2.1 void systemDefinition::initRandom (const int *N*, const int *rngSeed*)

Initialize a system of *N* atoms with random velocities and positions on a lattice. Net momeentum is automatically initialized to zero.

Parameters

in	<i>N</i>	Number of atoms to create
in	<i>rngSeed</i>	Random number generator seed

Definition at line 37 of file system.cpp.

References [atoms](#).

```

{
    float3 totMomentum;
    totMomentum.x = 0; totMomentum.y = 0; totMomentum.z = 0;

    if (N < 1) {
        throw customException ("N must be > 0");
        return;
    }

    srand(rngSeed);
    int chunk = OMP_CHUNK;
    unsigned int i;
    atoms.resize(N);
#pragma omp parallel private(i)
    {
#pragma omp for schedule(dynamic, chunk)
        for (i = 0; i < N; ++i) {
            if (i < N-1) {
                atoms[i].vel.x = (RNG-0.5);
                atoms[i].vel.y = (RNG-0.5);
                atoms[i].vel.z = (RNG-0.5);
                totMomentum.x += atoms[i].vel.x;
                totMomentum.y += atoms[i].vel.y;
                totMomentum.z += atoms[i].vel.z;
            } else {
                atoms[i].vel.x = -totMomentum.x;
                atoms[i].vel.y = -totMomentum.y;
            }
        }
    }
}

```

```

        atoms[i].vel.z = -totMomentum.z;
    }
    atoms[i].pos.x = (RNG)*box_.x;
    atoms[i].pos.y = (RNG)*box_.y;
    atoms[i].pos.z = (RNG)*box_.z;
    atoms[i].acc.x = 0;
    atoms[i].acc.y = 0;
    atoms[i].acc.z = 0;
}
}
}

```

4.9.2.2 void systemDefinition::initThermal (const int *N*, const float *Tset*, const int *rngSeed*, const float *dx*)

Initialize a system of *N* atoms with random velocities to meet a desired temperature and positions on a lattice. Net momentum is automatically initialized to zero.

Parameters

in	<i>N</i>	Number of atoms to create
in	<i>rngSeed</i>	Random number generator seed

Definition at line 83 of file system.cpp.

References atoms.

```

{
    float3 totMomentum;
    totMomentum.x = 0; totMomentum.y = 0; totMomentum.z = 0;

    if (N < 1) {
        throw customException ("N must be > 0");
        return;
    }

    srand(rngSeed);
    atoms.resize(N);

    // maxwell boltzmann distribution has mean 0 stdev kT/m in each dimension
    typedef std::tr1::linear_congruential<int, 16807, 0, (int)((1U << 31) -1 )
    > Myceng;
    Myceng eng;
    float sig = sqrt(Tset/mass_);
    std::tr1::normal_distribution<float> distribution(0.0,sig);
    float rannum = 0.0;
    float tmpT = 0.0;
    int index = 0;
    const int xs = floor(box_.x/dx);
    const int ys = floor(box_.y/dx);
    const int zs = floor(box_.z/dx);

    // initialize particle positions on a simple cubic lattice
    for (unsigned int x = 0; x < xs; ++x) {
        for (unsigned int y = 0; y < ys; ++y) {
            for (unsigned int z = 0; z < zs; ++z) {
                if (index < N) {
                    atoms[index].pos.x = x*dx;
                    atoms[index].pos.y = y*dx;
                    atoms[index].pos.z = z*dx;
                } else {
                    x = xs;
                    y = ys;
                    z = zs;
                    break;
                }
                index++;
            }
        }
    }

    for (unsigned int i = 0; i < N; ++i) {
        atoms[i].acc.x = 0;
        atoms[i].acc.y = 0;
        atoms[i].acc.z = 0;
        if (i < N-1) {
            atoms[i].vel.x = distribution(eng);
            atoms[i].vel.y = distribution(eng);
            atoms[i].vel.z = distribution(eng);
        }
    }
}

```



```

        totMomentum.x += atoms[i].vel.x;
        totMomentum.y += atoms[i].vel.y;
        totMomentum.z += atoms[i].vel.z;
    } else {
        atoms[i].vel.x = -totMomentum.x;
        atoms[i].vel.y = -totMomentum.y;
        atoms[i].vel.z = -totMomentum.z;
    }
    tmpT += (atoms[i].vel.x*atoms[i].vel.x + atoms[i].vel.y*
        atoms[i].vel.y + atoms[i].vel.z*atoms[i].vel.z);
}

// do velocity rescaling to get exactly the right T
tmpT *= mass_/(3.0*(N-1));
tmpT /= Tset;
tmpT = 1.0/tmpT;
for (unsigned int i = 0; i < N; i++) {
    atoms[i].vel.x = atoms[i].vel.x*sqrt(tmpT);
    atoms[i].vel.y = atoms[i].vel.y*sqrt(tmpT);
    atoms[i].vel.z = atoms[i].vel.z*sqrt(tmpT);
}
}

```

4.9.2.3 void systemDefinition::setPotential (pointFunction_t pp)

Assign pair potential function.

Sets the "host" pair potential function which also sets the GPU equivalent in [integrator.cu](#) if using CUDA.

Parameters

in	pp	Pointer to pair potential function
----	----	------------------------------------

Definition at line 21 of file system.cpp.

References potential.

```

{
/*#ifdef NVCC
    dev_potential = pp;
    cudaMemcpyFromSymbol(&potential, dev_potential, sizeof(pointFunction_t));
#else*/
    potential = pp;
//#endif
}

```

4.9.2.4 void systemDefinition::writeSnapshot ()

Print a snapshot of the system.

Write instantaneous snapshot of the system to a file called "trajectory.xyz" This file is appended not overwritten consecutively.

Definition at line 159 of file system.cpp.

References atoms.

```

{
static int snapNum = 0;
if (snapFile_ == NULL) {
    snapFile_ = fopen("trajectory.xyz", "w");
}

fprintf(snapFile_, "%d\nSnapshot #d\n", atoms.size(), snapNum);

for (unsigned int i = 0; i < atoms.size(); ++i) {
    fprintf(snapFile_, "%s\t%g\t%g\t%g\n", "A", atoms[i].pos.x, atoms
        [i].pos.y, atoms[i].pos.z);
}

snapNum++;
}

```

The documentation for this class was generated from the following files:

- /Users/nathanmahynski/CBEMDGPU/system.h
- /Users/nathanmahynski/CBEMDGPU/system.cpp

4.10 systemProps Class Reference

Holds all information about the CPU and GPU the simulation is being performed on.

```
#include <cudaHelper.h>
```

Public Member Functions

- **systemProps** (std::string fname)
- void [displayAllProps](#) ()
Display cpu host and device (if using GPUs) properties.
- `__host__` void [displayCudaProps](#) ()
Display the properties of the GPU being used.
- void [displayHost](#) ()
Display host name and other properties.
- int [maxThreadsPerBlock](#) (const int devID)
Returns the maximum number of threads per block on the GPU.
- int [maxGridDimX](#) (const int devID)
Returns the maximum number of blocks per grid on the GPU.
- int [numDevices](#) () const
Returns the number of GPUs found attached to the CPU.

4.10.1 Detailed Description

Holds all information about the CPU and GPU the simulation is being performed on.

Functions to assist in using CUDA and/or GPUs

Author

Nathan A. Mahynski

Date

11/21/13

Definition at line 26 of file cudaHelper.h.

4.10.2 Member Function Documentation

4.10.2.1 `__host__` void systemProps::displayCudaProps ()

Display the properties of the GPU being used.

Display properties of all CUDA capable devices

Definition at line 23 of file cudaHelper.cu.

Referenced by [displayAllProps](#)().

```

        {
checkCudaDevices_();

for (unsigned int i = 0; i < device_.size(); ++i) {
    *output_ << "# > Device Name: " << device_[i].name << endl;
    *output_ << "# -----" << endl;
    *output_ << "# Total Global Memory: " << device_[i].totalGlobalMem/1024
    << " KB" << endl;
    *output_ << "# Shared Memory available per Block: " << device_[i].
sharedMemPerBlock/1024 << " KB" << endl;
    *output_ << "# Registers per Thread Block: " << device_[i].regsPerBlock
    << endl;
    *output_ << "# Warp Size: " << device_[i].warpSize << endl;
    *output_ << "# Memory Pitch: " << device_[i].memPitch << endl;
    *output_ << "# Maximum Threads per Block: " << device_[i].
maxThreadsPerBlock << endl;
    *output_ << "# Maximum Thread Dimensions (Block): " << device_[i].
maxThreadsDim[0] << "x" << device_[i].maxThreadsDim[1] << "x" << device_[i].
maxThreadsDim[2] << endl;
    *output_ << "# Maximum Thread Dimensions (Grid): " << device_[i].
maxGridSize[0] << "x" << device_[i].maxGridSize[1] << "x" << device_[i].maxGridSize[2]
    << endl;
    *output_ << "# Total Constant Memory: " << device_[i].totalConstMem <<
    " bytes" << endl;
    *output_ << "# CUDA version: " << device_[i].major << "." << device_[i].
    .minor << endl;
    *output_ << "# Clock Rate: " << device_[i].clockRate << " kHz" << endl;
    *output_ << "# Texture Alignment: " << device_[i].textureAlignment << "
    bytes" << endl;
    if (device_[i].deviceOverlap == 0) {
        *output_ << "# Device Overlap: Not Allowed" << endl;
    } else {
        *output_ << "# Device Overlap: Allowed" << endl;
    }
    *output_ << "# Number of Multiprocessors: " << device_[i].
    multiProcessorCount << endl;
    *output_ << "# -----" << endl;
}
}


```

The documentation for this class was generated from the following files:

- /Users/nathanmahynski/CBEMDGPU/cudaHelper.h
- /Users/nathanmahynski/CBEMDGPU/cudaHelper.cu

4.11 SystemTest Class Reference

Inheritance diagram for SystemTest:



Protected Member Functions

- virtual void **SetUp** ()

Protected Attributes

- [systemDefinition](#) a
- int **natoms**
- float **mass**
- float **L**
- float **eps**

- float **sigma**
- [nvt_NH](#) **integrate**

4.11.1 Detailed Description

Definition at line 11 of file unittests.cpp.

The documentation for this class was generated from the following file:

- /Users/nathanmahynski/CBEMDGPU/unittests.cpp